

CPSC 213 – Assignment 5

Memory Management

Due: Monday, October 19, 2020 at 11:59pm

After an 8-hour, no-penalty grace period, no late assignments will be accepted.

Learning Objectives

Here are the learning objectives of this assignment, which will be examined in the corresponding quiz. They are a subset of the unit learning objectives listed on Slide 2 of Unit 1c.

After completing this assignment you should be able to:

1. identify and correct dangling-pointer and memory leak bugs in C caused by improper use of `free()`;
2. write C code that uses techniques that avoid dynamic allocation as a way to minimize memory-allocation bugs; and
3. write C code that uses reference counting as a way to minimize memory-allocation bugs.

Goal

In this week's assignment you will do two things. First, you get additional practice reading assembly code, figuring out what C program it came from, and counting memory references — skills that will be useful on the midterm and beyond.

Second, we will examine dynamic allocation and de-allocation in C. The goal here is to help you clarify your understanding of this topic, particularly dangling pointers, memory leaks, and how to avoid them. You'll start by implementing a simple memory allocator, to get you more familiar with how it works under the hood. Then, you will examine a program that contains a dangling pointer bug to identify the problem and fix it.

Download the code file

Download the file www.students.cs.ubc.ca/~cs-213/cur/assignments/a5/code.zip. It contains the following files, some in subdirectories named `q2` and `q3`:

- | | |
|-------------------------------|------------------|
| 1. <code>q1.[cs]</code> | – Question 1 |
| 2. <code>refcount.[ch]</code> | – Questions 2, 3 |

- 3. `q2: main.c, mymalloc.[ch], Makefile` – Question 2
- 4. `q3: {list,tree,element}.[ch], main.c, and Makefile` – Question 3

Question 1 — Reading Assembly Code [10%]

Read the included `q1.c` and `q1.s` files. You will see that `q1.c` declares two structs, allocates some objects from those structs and initializes them. It has a procedure named `q1()` that is blank. The code listed in `q1.s` implements the code in this procedure (without the procedure call itself). Read `q1.s` carefully and run it through the simulator to figure out how it manipulates these structs. Then do the following.

1. [2%] Comment every code line of `q1.s` with a high-level comment (i.e, a C-like comment).
2. [6%] Modify the procedure `q1()` in `q1.c` to perform the same computation as `q1.s`. Do not modify the other parts of this C file; they are used to test and mark your code. Note that you can optionally provide different values for the structs on the command line. The C code you write must work for arbitrary values.
3. [2%] Examine the code you wrote for `q1()`. Count the *minimum* number of memory reads and writes that are required to execute these five statements together. Note these numbers may be different from `q1.s`, which takes each line individually (e.g., reading the value of `i` each time, which is not really necessary). Ignore register allocation for this questions; i.e., you can assume that you have an infinite number of registers available.

Place these two numbers in the file `q1.txt` on two separate lines with the number of reads listed first (just these numbers and nothing else). Then carefully explain your answer by listing each of the reads and writes that are required (describe the access using variable names) in the file `q1-desc.txt`.

Question 2 — malloc and free [55%]

In this assignment, you will implement a functioning but simple version of `malloc` and `free` from scratch. Real implementations of `malloc/free` must deal with two issues that we will ignore: fragmentation and efficient allocation, which makes the real implementation much more complex (over 5000 lines of code) than what you will write. Your goal is simply to ensure that (a) every active allocation is located in its own, distinct location in memory and that (b) memory that is freed is reused whenever possible in subsequent calls to `malloc`.

The goal is to arrive at a simple, yet functional implementation which is able to service requests using an *explicit free list* (textbook section 9.9.13; basically, freed blocks are added into a linked list and searched during `malloc`). You will be guided through the approach, but start early because it's easy to make mistakes and have to spend time debugging.

To start off, look at the code sample provided. It has a Makefile and two main modules: `mymalloc.c`, `{c,h}` and `main.c`. You will modify and hand in the implementation in `mymalloc.c`, but **leave**

mymalloc.h and main.c alone (you can modify them for local testing if you like, but the autograder will not use your versions). You can add helper functions, variables, and constants as you wish in `mymalloc.c`.

`mymalloc` exports three functions. Read through the documentation in the header file to know what they do, and see `main.c` for how they get used. The initial implementation you are given does not reuse freed memory. Instead, every call to `malloc` takes its memory for the current top of the heap, and then advances the heap top by the size of the allocation. This is simple and fast, but of course means that every program has a memory leak. Your goal is to change `malloc` and `free` so that freed memory is reused whenever possible.

`main.c` implements three test-cases which will stress your `malloc` implementation. Your goal is to pass all of the tests. The autograder will use the exact same tests, so if you pass every test locally you should be able to pass the tests on the autograder. To compile, run `make` in this directory to make the `mymalloc` binary.

To run a test, run `./mymalloc <test name>`, or use `./mymalloc all` to run every test in sequence. The provided implementation of `malloc` passes `sanity`, but fails the others.

Implementing malloc and free

1. The first problem you need to solve is to record the *size* of each allocated block as part of the block, in a prefix *metadata* hidden from the caller. You'll need to know this size when returning blocks to the *freelist* when `myfree` is called. You'll use the same technique here as we do to store the reference count for reference counted objects, so use the provided `refcount.c` as a guide.

Modify your implementation to add the *metadata prefix* to the beginning of the allocated block, thus increasing its size, and to shift the return value pointer to bypass this prefix. Be sure that the address you return is properly aligned to an 8-byte boundary, which is something `malloc` must always do.

Test your implementation by (a) ensuring that it still passes the `sanity` test and (b) adding a test to check that the size you store in the object has the correct value (e.g., by adding that test to the `test_sanity` procedure used in the `sanity` test).

2. Next, you'll implement `myfree` by building a *singly-linked list* out of the blocks of free memory; each call to `myfree` adding to the top of the list. You can take a look at `list.c` and `list.h` from Question 3 for an example implementation. Note, however, that `list.c` implements a *doubly-linked list*, so your implementation will be simpler as you need only a `next` pointer and no `prev` (or `elem`) pointer. Note also, that you will store this `next` pointer inside of the freed block itself.

Do the following. Create a static `head` pointer to point to the most recently freed block of memory. On each call to `myfree`, store the current value of the `head` pointer into the block being freed, as its `next` pointer, and change the `head` to point to this freed block; thus forming a linked list out of the freed blocks.

Test your implementation to make sure it still passes the `test_sanity` test and then temporarily modify `myfree` to check that the freelist has the correct size after each call to `myfree`. Note that at this point `mymalloc` is not using the freelist and so the size of the freelist is equal to the number

of times `myfree` has been called, which you can determine using a static counter variable. Once this test passes, comment out this test code before proceeding to Step 3, because once `mymalloc` starts using the freelist, your test code will no longer be valid. *Note: pointers on the autograding machine will be 8 bytes in size.*

3. Finally, change `mymalloc` to use the free list. Check the *size metadata* of each block on the freelist to find one that is big enough to hold the allocation; it is fine to use a block that is bigger than the requested size. You can experiment with different ways to choose a matching block: possible choices include *first-fit*, which returns the first block found that is big enough, or *best-fit*, which returns the smallest suitable block. Make sure to unlink the returned block from the free list.
4. By this point, you should be able to pass all the tests. If you don't (which is extremely likely at first), you'll have to debug your implementation by adding `printf` statements and/or using `gdb` or `lldb`. Start your testing with the `freelist1` test, which is the simplest test and thus the easiest to use for debugging. Don't try the other tests until this test passes. Then do the same with each of the other tests in this order: `freelist2`, `fixedsize`, and `random`.

Hand in only the file `mymalloc.c`.

Question 3 – Reference Counting [35%]

NOTE: This question has been moved to Assignment 6. Do not submit anything for this question this week, but of course you might start on it early if you have time. The weighting of the other two questions will be adjusted to 15% and 85%.

The program in the `q3` directory can be built using `make` while in that directory. You can then run the program by typing

```
./main A B C D
```

The program takes an arbitrary number of string arguments. It is *supposed* to print out a binary search tree, which is formed out of a subset of arguments, twice. However, as you can clearly see, the program does not produce this output; it correctly prints only one tree. **It has a bug.** To see the correct output of the program, you can remove all the `free()` calls in the program, replacing the dangling pointer bug with a memory leak.

The implementation uses two libraries, `list` and `tree`, implementing a doubly-linked list and a binary search tree respectively. Both libraries are free of bugs if used independently, but exhibit a bug when used together.

1. Find the bug and write a careful description of it in the file `q3.txt`. The bug is a dangling pointer somewhere. You may also want to read the header files carefully to see whether the libraries are being misused.
2. Now, fix the bug using reference counting. You must use the provided `refcount.c` and `.h` files for reference counting (the Makefile is already setup to use them; add `"#include \"refcount.h\""` to C files as needed). You can change the interface in `element.h` and any of the implementation in `element.c` (retaining its ability to store strings and numbers). You should

not change any of the other `.h` files, and make minimal changes to the other `.c` files aside from adding reference counting (e.g. calls to element-specific `keep_ref/free_ref` functions that you define). Elements must still be shared between the list and tree - you should not simply make copies of the elements (which would introduce inefficiency).

3. After fixing the bug, test your program to ensure that it produces the correct output and that `valgrind` reports no memory leaks. To do so, you must run the program through `valgrind`, which is available on the undergrad servers but likely not on other platforms you might use for other parts of the assignment. The test you need to run is the following:

```
valgrind --tool=memcheck --leak-check=yes ./main <arguments>
```

If you have a memory leak you will see an error that looks something like this:

```
==31419== LEAK SUMMARY:
==31419==      definitely lost: 240 bytes in 2 blocks
==31419==      indirectly lost: 240 bytes in 2 blocks
```

If you have a dangling pointer you may see an error that looks something like this:

```
==31993== Invalid read of size 1
==31993==      at 0x4E7D000: vfprintf (vfprintf.c:1629)
```

If you eliminate both bugs (your goal) then you should see output something like this:

```
==1272== HEAP SUMMARY:
==1272==      in use at exit: 0 bytes in 0 blocks
==1272==    total heap usage: 4 allocs, 4 frees, 480 bytes allocated
==1272==
==1272== All heap blocks were freed -- no leaks are possible
==1272==
==1272== For counts of detected and suppressed errors, rerun with: -v
==1272== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 2 from 2)
```

Or you might see (depending on which version of `valgrind` you use):

```
==9402== LEAK SUMMARY:
==9402==      definitely lost: 0 bytes in 0 blocks
==9402==      indirectly lost: 0 bytes in 0 blocks
==9402==      possibly lost: 0 bytes in 0 blocks
==9402==    still reachable: 4,096 bytes in 1 blocks
==9402==         suppressed: 25,084 bytes in 373 blocks
```

What to Hand In

Use the `handin` program.

The assignment directory is `~/cs-213/a5`. It should contain the following *plain-text* files:

1. The files `q1.s`, `q1.c`, `q1.txt`, and `q1-desc.txt` for Question 1;
2. a subdirectory named `q2` just like you see in the code file;
3. in the `q2` directory: `mymalloc.c`;